


Pointers, Arrays, Memory: AKA the cause of those F@#)(#@*(Segfaults




Announcements!

- HW1 is due Friday, January 28
- Lab 1 is due Monday, January 31
- Discussion and lab schedule has been uploaded to the website
- Office hours schedule coming soon
- Project 1 estimated release... Wednesday
 - Trying for maximum debugging *and debuggability* for snek!


Address vs. Value

- Consider memory to be a ***single*** huge array
 - Each cell of the array has an address associated with it
 - Each cell also stores some value
 - For addresses do we use signed or unsigned numbers? Negative address?!
 - Answer: Addresses are ***unsigned***
- Don't confuse the address referring to a memory location with the value stored there



Pointers

- An *address* refers to a particular memory location; e.g., it points to a memory location
- *Pointer*: A variable that contains the address of a variable



Types of Pointers

- Pointers are used to point to any kind of data (`int`, `char`, a `struct`, a pointer to a pointer to a pointer to a `char`, etc.)
- Normally a pointer only points to one type (`int`, `char`, a `struct`, etc.).
 - `void *` is a type that can point to anything (generic pointer)
 - Use `void *` sparingly to help avoid program bugs, and security issues, and other bad things!
 - Can convert types (BUT BE CAREFUL):

```
void *a = ....
int *p = (int *) a;    /* p now points to the same place as a,
                        but is treated as a pointer to an int */
int **q = (int **) a; /* q now points to the same place as a,
                        but is treated as a pointer to a pointer to an int */
```
- You can even have pointers to functions...
 - `int (*fn) (void *, void *) = &foo`
 - `fn` is a function that accepts two `void *` pointers and returns an `int` and is initially pointing to the function `foo`.
 - `(*fn) (x, y)` will then call the function

NULL pointers...

- The pointer of all 0s is special
 - The "NULL" pointer, like in Java, python, etc...
- If you write to or read a null pointer, your program should crash immediately
 - The memory is set up so that this should never be valid
- Since "0 is false", its very easy to do tests for null:
 - `if(!p) { /* p is a null pointer */ }`
 - `if(q) { /* q is not a null pointer */ }`

More C Pointer Dangers

- *Declaring a pointer just allocates space to hold the pointer – it does not allocate the thing being pointed to!*
- Local variables in C are not initialized, they may contain anything (aka “garbage”)
- What does the following code do?

```
void f()  
{  
    int *ptr;  
    *ptr = 5;  
}
```

Pointers and Structures

```
typedef struct {  
    int x;  
    int y;  
} Point;
```

```
Point p1;  
Point p2;  
Point *paddr;  
paddr = &p2;
```

```
/* dot notation */  
int h = p1.x;  
p2.y = p1.y;
```

```
/* arrow notation */  
int h = paddr->x;  
int h = (*paddr).x;
```

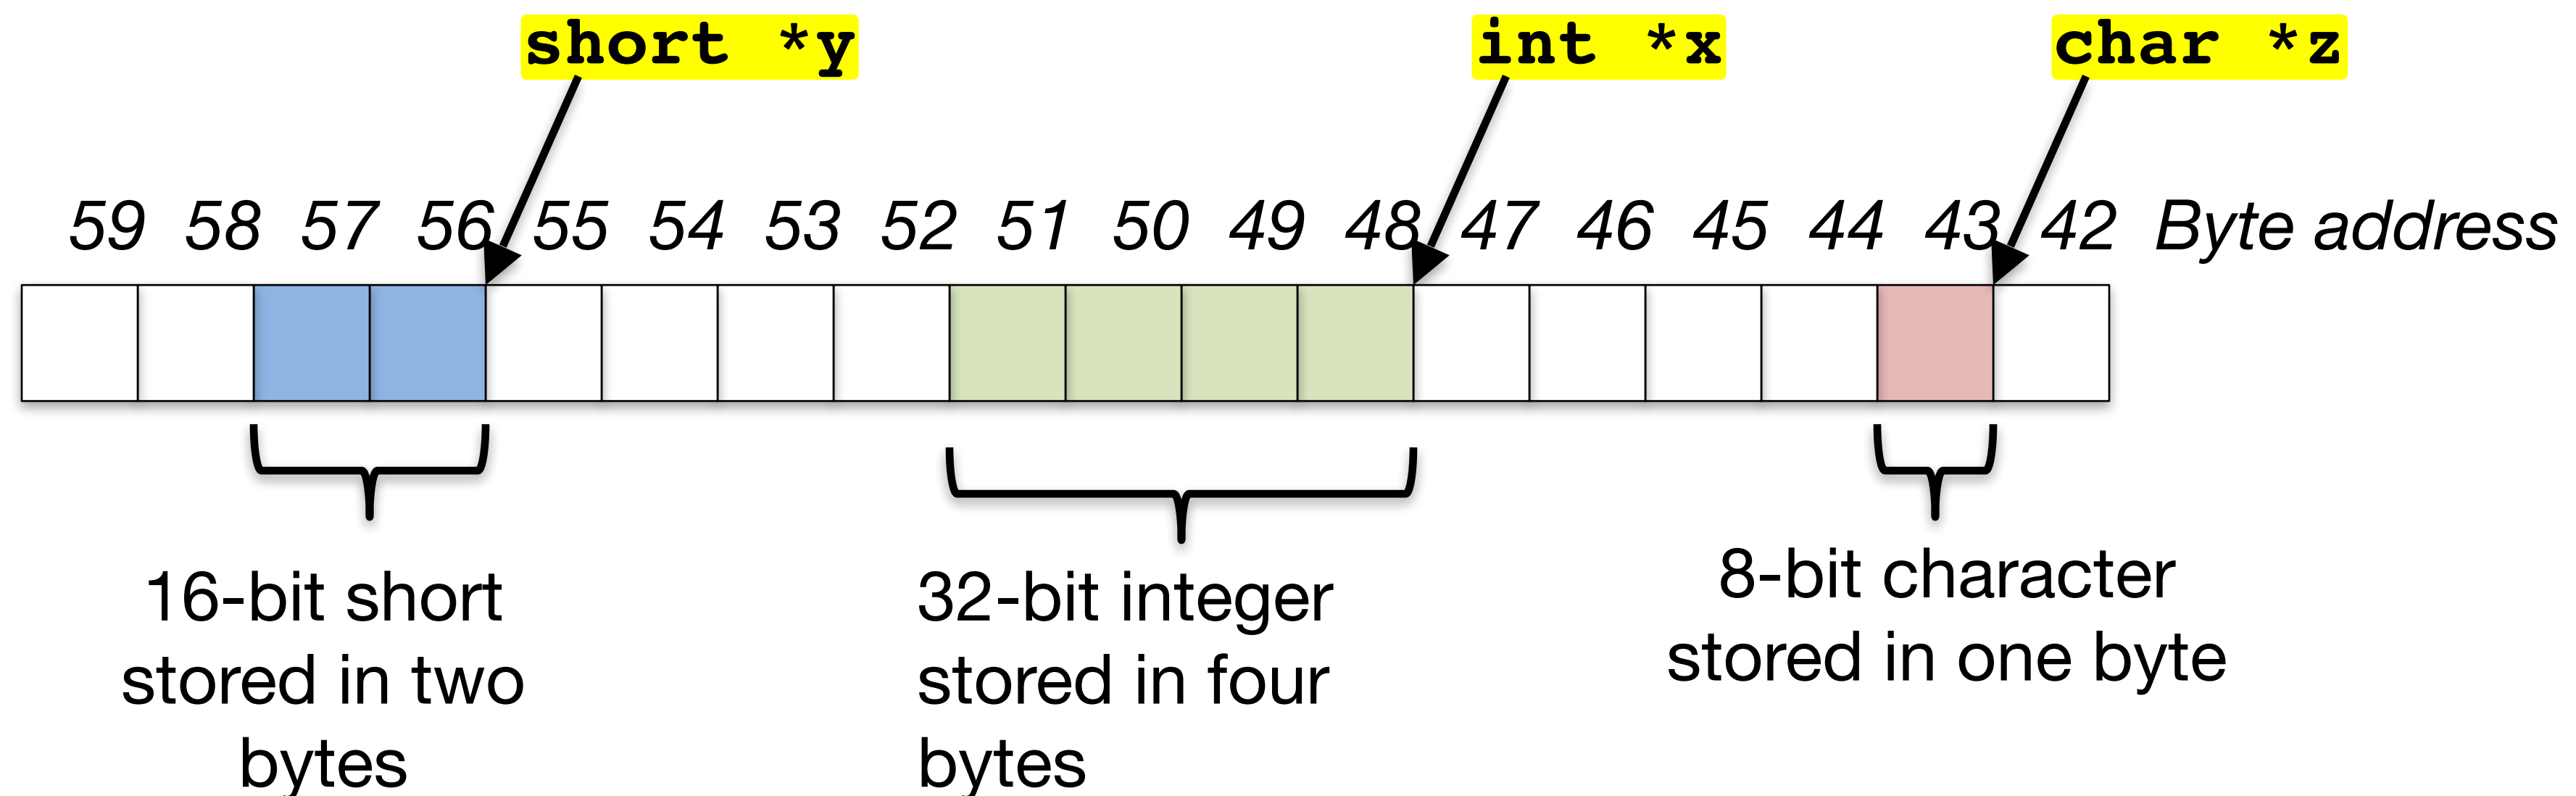
```
/* This works too:  
   copies all of p2 */  
p1 = p2;  
p1 = *paddr;
```


Pointers in C

- Why use pointers?
 - If we want to pass a large struct or array, it's easier / faster / etc. to pass a pointer than the whole thing
 - Otherwise we'd need to copy a huge amount of data
 - You notice in Java that more complex objects are passed by reference....
Under the hood this is a pointer
 - In general, pointers allow cleaner, more compact code
- So what are the drawbacks?
 - Pointers are probably the single largest source of bugs in C, so be careful anytime you deal with them
 - Most problematic with dynamic memory management—coming up next time
 - Dangling references and memory leaks

Pointing to Different Size Objects

- Modern machines are “byte-addressable”
 - Hardware’s memory composed of 8-bit storage cells, each has a unique address
- A C pointer is just abstracted memory address
- Type declaration tells compiler how many bytes to fetch on each access through pointer
 - E.g., 32-bit integer stored in 4 consecutive 8-bit bytes
- But we actually want “word alignment”
 - Some processors will not allow you to address 32b values without being on 4 byte boundaries
 - Others will just be very slow if you try to access “unaligned” memory.



sizeof() operator

- **sizeof (type)** returns number of bytes in object
 - But number of bits in a byte is not standardized technically
 - In olden times, when dragons roamed the earth, bytes could be 5, 6, 7, 9 bits long
 - Includes any padding needed for alignment
 - So that every `int` will start at a boundary divisible by 4...
- By Standard C99 definition, **sizeof (char) == 1**
- Can take **sizeof (arg)**, or **sizeof (structtype)**
- We'll see more of sizeof when we look at dynamic memory management
- **sizeof is *not* a function!** It is a compile-time operation

Pointer Arithmetic

pointer + number

e.g., *pointer + 1*

```
char *p;  
char a;  
char b;  
  
p = &a;  
p += 1;
```

pointer - number

adds 1 something to a pointer

```
int *p;  
int a;  
int b;  
  
p = &a;  
p += 1;
```

In each, p now points to b
(Assuming compiler doesn't
reorder variables in memory.)

Never code like this!!!!

Adds `1*sizeof(char)`
to the memory address

Adds `1*sizeof(int)`
to the memory address

Pointer arithmetic should be used cautiously

Basic rule for pointer arithmetic


- We cover it for two reasons
 - You may encounter code using this in the future...
 - You need to understand this to understand how this code gets converted to assembly
- Look at the type the pointer points **to**
 - So a (char *) points to a (char), while a (char **) points to a (char *)
 - The **actual value used by the compiler (NOT THE VALUE YOU USE)** is the size of what you are pointing to time the amount to increment
- So **under the hood**: `char *c; char **d;`
- `(c + 5) -> c + sizeof(char) * 5 -> c + 5`
- `(d + 7) -> d + sizeof(char *) * 5 -> d + 20`

Changing a Pointer Argument?

- What if want function to change a pointer?
- What gets printed?

```
void inc_ptr(int *p)
{
    p = p + 1;
}

int A[3] = {50, 60, 70};
int* q = A;
inc_ptr(q);
printf("*q = %d\n", *q);
```




Pointer to a Pointer

- Solution! Pass a pointer to a pointer, declared as ****h**
- Now what gets printed?

```
void inc_ptr(int **h)
{   *h = *h + 1;   }

int A[3] = {50, 60, 70};
int* q = A;
inc_ptr(&q);
printf("*q = %d\n", *q);
```



It can never end...

- You can have something like this:
`int *****x;`
- x is a pointer to a pointer to a pointer to a pointer to a pointer to a pointer to a pointer to a pointer to a pointer to an integer!

Conclusion on Pointers...

- All data is in memory
 - Each memory location has an address to use to refer to it and a value stored in it
- Pointer is a C version (abstraction) of a data address
 - * “follows” a pointer to its value
 - & gets the address of a value
- C is an efficient language, but leaves safety to the programmer
 - Variables not automatically initialized
 - Use pointers with care: they are a common source of bugs in programs

Structures Revisited

- A "struct" is really just an instruction to C on how to arrange a bunch of bytes in a bucket...
- ```
struct foo {
 int a;
 char b;
 struct foo *c;
}
```
- Provides enough space and ***aligns*** the data with padding  
So actual layout on a 32b architecture will be:
  - 4-bytes for A
  - 1 byte for b
  - 3 unused bytes
  - 4 bytes for C
- `sizeof(struct foo) == 12`



# Plus also Unions

- A "union" is also instruction to C on how to arrange a bunch of bytes
- ```
union foo {  
    int a;  
    char b;  
    union foo *c;  
}
```
- Provides enough space for the *largest element*
- ```
union foo f;
f.a = 0xDEADB33F; /* treat f as an integer and store
 that value */
f.c = &f; /* treat f as a pointer of type
 "union foo *" and store the
 address of f in itself */
```

# C Arrays

- Declaration:

```
int ar[2];
```

declares a 2-element integer array: just a block of memory which is uninitialized. The number of elements is static in the declaration, you can't do "**int ar[x]**" where x is a variable

```
int ar[] = {795, 635};
```

declares and initializes a 2-element integer array

# Array Name / Pointer Duality

- *Key Concept:* Array variable is simply a “pointer” to the first (0th) element
- So, array variables are ***almost*** identical to pointers
  - `char *string` and `char string[]` are nearly identical declarations
    - Differ in subtle ways: incrementing & declaration of filled arrays
- Consequences:
  - `ar[32]` is an array variable with 32 elements, but works like a pointer
  - `ar[0]` is the same as `*ar`
  - `ar[2]` is the same as `*(ar+2)`
  - Can use pointer arithmetic to access arrays

# Arrays and Pointers

- Array  $\approx$  pointer to the initial element
  - $a[i] \equiv *(a+i)$
- An array is passed to a function as a pointer
  - The array size is **lost**!
- Usually bad style to interchange arrays and pointers
  - Avoid pointer arithmetic!
    - Especially avoid things like `ar++`;

## Passing arrays:

Really `int *array`      Must explicitly pass the size

```
int
foo(int array[],
 unsigned int size)
{
 ... array[size - 1] ...
}

int
main(void)
{
 int a[10], b[5];
 ... foo(a, 10)... foo(b, 5) ...
}
```

# C Arrays are Very Primitive

- An array in C does not know its own length, ***and its bounds are not checked!***
  - Consequence: We can accidentally ***access off the end of an array***
  - Consequence: We must pass the array ***and its size*** to any procedure that is going to manipulate it
- Segmentation faults and bus errors:
  - These are VERY difficult to find;  
be careful! (You'll learn how to debug these in lab)
  - But also “fun” to exploit:
    - “Stack overflow exploit”, maliciously write off the end of an array on the stack
    - “Heap overflow exploit”, maliciously write off the end of an array on the heap



# C Strings

- String in C is just an array of characters

```
char string[] = "abc";
```

- How do you tell how long a string is?
  - Last character is followed by a 0 byte (aka “null terminator”):  
written as 0 (the number) or '\0' as a character
  - Important danger: string length operation does **not** include the null terminator when you ask for length of a string!

```
int strlen(char s[])
{
 int n = 0;
 while (s[n] != 0) {
 n++;
 }
 return n;
}

int strlen(char s[])
{
 int n = 0;
 while (*(s++) != 0) {
 n++;
 }
 return n;
}
```

# Use Defined Constants

- Array size  $n$ ; want to access from  $0$  to  $n-1$ , so you should use counter AND utilize a variable for declaration & incrementation
  - Bad pattern

```
int i, ar[10];
for(i = 0; i < 10; i++){ ... }
```
  - Better pattern

```
const int ARRAY_SIZE = 10;
int i, a[ARRAY_SIZE];
for(i = 0; i < ARRAY_SIZE; i++){ ... }
```
- ***SINGLE SOURCE OF TRUTH***
  - You're utilizing indirection and avoiding maintaining two copies of the number 10
  - DRY: "Don't Repeat Yourself"
  - And don't forget the  $<$  rather than  $<=$ :  
When Nick took 60c, he lost a day to a "segfault in a malloc called by printf on large inputs":  
Had a  $<=$  rather than a  $<$  in a single array initialization!

# Arrays and Pointers

```
int
foo(int array[],
 unsigned int size)
{
 ...
 printf("%d\n", sizeof(array));
}

int
main(void)
{
 int a[10], b[5];
 ... foo(a, 10)... foo(b, 5) ...
 printf("%d\n", sizeof(a));
}
```

What does this print? **4**

... because `array` is really a pointer (and a pointer is architecture dependent, but likely to be 4 or 8 on modern 32-64 bit machines!)

What does this print? **40**

# Arrays and Pointers

```
int i;
int array[10];

for (i = 0; i < 10; i++)
{
 array[i] = ...;
}
```

```
int *p;
int array[10];

for (p = array; p < &array[10]; p++)
{
 *p = ...;
}
```

These code sequences have the same effect!

But the former is ***much more readable***:

Especially don't want to see code like `ar++`

# Arrays And Structures And Pointers

- ```
typedef struct bar {  
    char *a;      /* A pointer to a character */  
    char b[18];   /* A statically sized array  
                  of characters */  
} Bar;  
  
...  
Bar *b = (Bar*) malloc(sizeof(struct bar));  
b->a = malloc(sizeof(char) * 24);
```
- Will require 24 bytes on a 32b architecture for the structure:
 - 4 bytes for a (its a pointer)
 - 18 bytes for b (it is 18 characters)
 - 2 bytes padding (needed to align things)

Some Code Examples

- **`b->b[5] = 'd'`**
 - Location written to is 10th byte pointed to by b...
`*((char *) b + 4 + 5) = 'd'`
- **`b->a[5] = 'c'`**
 - location written to is the first word pointed to by b, treat that as a pointer, add 5, and write 'c' there...
aka `*(* (char **) b) + 5) = 'c'`
- **`b->a = b->b`**
 - Location written to is the first word pointed to by b
 - Value it is set to is b's address + 4)...
aka `* ((char **)b) = ((char *) b) + 4`

When Arrays Go Bad: Heartbleed

- In TLS encryption, messages have a length...
 - And get copied into memory before being processed
- One message was “Echo Me back the following data, its this long...”
 - But the (different) echo length wasn’t checked to make sure it wasn’t too big...

```
M 5 HB L=5000 107:Ou17;GET / HTTP/1.1\r\n
Host: www.mydomain.com\r\nCookie: login=1
17kf9012oeu\r\nUser-Agent: Mozilla...
```

- So you send a small request that says “read back a lot of data”
 - And thus get web requests with auth cookies and other bits of data from random bits of memory...

Concise `strlen()`

```
int strlen(char *s)
{
    char *p = s;
    while (*p++)
        ; /* Null body of while */
    return (p - s - 1);
}
```

What happens if there is no zero character at end of string?

Arguments in `main ()`

- To get arguments to the main function, use:
 - `int main(int argc, char *argv[])`
- What does this mean?
 - `argc` contains the number of strings on the command line (the executable counts as one, plus one for each argument). Here `argc` is 2:
 - `unix% sort myFile`
 - `argv` is a pointer to an array containing the arguments as strings
 - Since it is an array of pointers to character arrays
 - Sometimes written as `char **argv`

Example

- `foo hello 87 "bar baz"`
- `argc = 4 /* number arguments */`
- `argv[0] = "foo",`
`argv[1] = "hello",`
`argv[2] = "87",`
`argv[3] = "bar baz",`
- Array of pointers to strings

Endianness...

- Consider the following
- ```
union confuzzle { int a; char b[4]; };
union confuzzle foo;
foo.a = 0x12345678;
```
- In a 32b architecture, what would foo.b[0] be?  
0x12? 0x78?
- Its actually dependent on the architecture's "endianness"
  - Big endian: The first character is the most significant byte: 0x12
  - Little endian: The first character is the least significant byte: 0x78

# Endianness and You...

- It generally doesn't matter if you write portable C code running on one computer...
  - After all, you shouldn't be treating an integer as a series of raw bytes
  - Well, it matters when you take CS161:  
x86 is little endian and you may write an address as a string
- It does matter when you want to communicate across computers...
  - The "network byte order" is big-endian, but your computer is likely to be little-endian
    - x86, RISC-V, Apple M1 in practice are all little-endian
- Endian conversion functions:
  - `ntohs()`, `htons()`: Convert 16 bit values from your native architecture to network byte order and vice versa
  - `ntohl()`, `htonl()`: Convert 32 bit values from your native architecture to network byte order and vice versa

# C Memory Management


- How does the C compiler determine where to put all the variables in machine's memory?
- How to create dynamically sized objects?
- To simplify discussion, we assume *one program runs at a time*, with access to all of memory.
- Later, we'll discuss ***virtual memory***, which lets multiple programs all run at same time, each thinking they own all of memory
  - The only real addition is the C runtime has to say "Hey operating system, gimme a big block of memory" when it needs more memory

# C Memory Management

- Program's address space contains (32 bits assumed here)  
4 regions:  
Memory Address  
~ **FFFF FFFF**<sub>hex</sub>

- **stack**: local variables inside functions, grows downward
- **heap**: space requested for dynamic data via `malloc()` resizes dynamically, grows upward
- **static data**: variables declared outside functions, does not grow or shrink. Loaded when program starts, can be modified.
- **code**: loaded when program starts, does not change
- 0x0000 0000 hunk is reserved and unwriteable/unreadable so you crash on null pointer access

~ **0000 0000**<sub>hex</sub>





# Where are Variables Allocated?

- If declared outside a function, allocated in “static” storage
- If declared inside function, allocated on the “stack” and freed when function returns
  - `main()` is treated like a function
- For both of these types of memory, the management is automatic:
  - You don't need to worry about deallocating when you are no longer using them
  - But a variable ***does not exist anymore*** once a function ends!  
Big difference from Java

```
int myGlobal;
main() {
 int myTemp;
}
```




# The Stack

- Every time a function is called, a new "stack frame" is allocated on the stack
- Stack frame includes:
  - Return address (who called me?)
  - Arguments
  - Space for local variables
- Stack frames uses contiguous blocks of memory; stack pointer indicates start of stack frame
- When function ends, stack pointer moves up; frees memory for future stack frames
- We'll cover details later for RISC-V processor

```
fooA() { fooB(); }
fooB() { fooC(); }
fooC() { fooD(); }
```

**Stack Pointer** →



# Stack Animation

- Last In, First Out (LIFO) data structure


```
main ()
{ a(0);
}

void a (int m)
{ b(1);
}

void b (int n)
{ c(2);
}

void c (int o)
{ d(3);
}

void d (int p)
{
}
```



# Managing the Heap

C supports functions for heap management:

- **malloc()** allocate a block of ***uninitialized*** memory
- **calloc()** allocate a block of ***zeroed*** memory
- **free()** free previously allocated block of memory
- **realloc()** change size of previously allocated block
  - careful – it might move!
    - And it ***will not update other pointers pointing to the same block of memory***

# Malloc()

- **`void *malloc(size_t n):`**
  - Allocate a block of uninitialized memory
  - NOTE: Subsequent calls probably will not yield adjacent blocks
  - `n` is an integer, indicating size of requested memory block in bytes
  - `size_t` is an unsigned integer type big enough to “count” memory bytes
  - Returns `void*` pointer to block; `NULL` return indicates no more memory (check for it!)
  - Additional control information (including size) stored in the heap for each allocated block.
- Examples:
  - *“Cast” operation, changes type of a variable.  
Here changes (`void *`) to (`int *`)*  
`int *ip;`  
`ip = (int *) malloc(sizeof(int));`
  - `typedef struct { ... } TreeNode;`  
`TreeNode *tp = (TreeNode *) malloc(sizeof(TreeNode));`
- `sizeof` returns size of given type in bytes, ***necessary if you want portable code!***

# And then free()

- **void free(void \*p) :**
  - **p** is a pointer containing the address originally returned by **malloc()**
- Examples:
  - ```
int *ip;  
ip = (int *) malloc(sizeof(int)) ;  
... ..  
free((void*) ip) ; /* Can you free(ip) after ip++ ? */
```
 - ```
typedef struct {... } TreeNode;
TreeNode *tp = (TreeNode *) malloc(sizeof(TreeNode)) ;
... ..
free((void *) tp) ;
```
- When you free memory, you must be sure that you pass the original address returned from **malloc()** to **free()**; Otherwise, crash (or worse)!



# Using Dynamic Memory


```
typedef struct node {
 int key;
 struct node *left; struct node
 *right;
} Node;

Node *root = NULL;

Node *create_node(int key, Node *left,
 Node *right){
 Node *np;
 if(!(np =
 (Node*) malloc(sizeof(Node)))){
 printf("Memory exhausted!\n");
 exit(1);}
 else{
 np->key = key;
 np->left = left;
 np->right = right;
 return np;
 }
}
```

```
void insert(int key, Node **tree){
 if ((*tree) == NULL){
 (*tree) = create_node(key, NULL,
 NULL);
 }
 else if (key <= (*tree)->key){
 insert(key, &((*tree)->left));
 }
 else{
 insert(key, &((*tree)->right));
 }
}
```

```
int main(){
 insert(10, &root);
 insert(16, &root);
 insert(5, &root);
 insert(11, &root);
 return 0;
}
```



# Observations

- Code, Static storage are easy: they never grow or shrink
- Stack space is relatively easy: stack frames are created and destroyed in last-in, first-out (LIFO) order
- Managing the heap is tricky: memory can be allocated / deallocated at any time
  - If you forget to deallocate memory: “Memory Leak”
    - Your program ***will eventually run out of memory***
  - If you call free twice on the same memory: “Double Free”
    - Possible ***crash or exploitable vulnerability***
  - If you use data after calling free: “Use after free”
    - Possible ***crash or exploitable vulnerability***

# When Memory Goes Bad...

## Failure To Free

- #1: Failure to free allocated memory
  - "memory leak"
- Initial symptoms: nothing
  - Until you hit a critical point, memory leaks aren't actually a problem
- Later symptoms: performance drops off a cliff...
  - Memory hierarchy behavior tends to be good just up until the moment it isn't...
    - There are actually a couple of cliffs that will hit
- And then your program is killed off!
  - Because the OS goes "Nah, not gonna do it" when you ask for more memory

# When Memory Goes Bad: Writing off the end of arrays...

- EG...
  - ```
int *foo = (int *) malloc(sizeof(int) * 100);  
int i;  
....  
for(i = 0; i <= 100; ++i) {  
    foo[i] = 0;  
}
```
- Corrupts other parts of the program...
 - Including internal C data used by `malloc()`
- May cause crashes later

When Memory Goes Bad: Returning Pointers into the Stack

- It is OK to pass a pointer to stack space down
 - EG:

```
char [40] foo;  
int bar;  
...  
strncpy(foo, "102010", strlen("102010")+1);  
baz(&bar);
```
- It is catastrophically bad to return a pointer to something in the stack...
 - EG

```
char [50] foo;  
....  
return foo;
```
- The memory will be overwritten when other functions are called!
 - So your data no longer exists... And writes can overwrite key pointers causing crashes!

When Memory Goes Bad: Use After Free

- When you keep using a pointer..
 - `struct foo *f`
 ...
 `f = malloc(sizeof(struct foo)) ;`
 ...
 `free(f)`
 ...
 `bar(f->a) ;`
- Reads after the free may be corrupted
 - As something else takes over that memory. Your program will probably get wrong info!
- Writes ***corrupt*** other data!
 - Uh oh... Your program crashes later!

When Memory Goes Bad: Forgetting Realloc Can Move Data...

- When you realloc it can copy data...
 - `struct foo *f = malloc(sizeof(struct foo) * 10);`
...
`struct foo *g = f;`
....
`f = realloc(sizeof(struct foo) * 20);`
- Result is g *may* now point to invalid memory
 - So reads may be corrupted and writes may corrupt other pieces of memory

When Memory Goes Bad: Freeing the Wrong Stuff...

- If you `free()` something never `malloc'ed()`
 - Including things like

```
struct foo *f = malloc(sizeof(struct foo) * 10)
...
f++;
...
free(f)
```
- Malloc/free may get confused..
 - Corrupt its internal storage or erase other data...

When Memory Goes Bad: Double-Free...

- EG...
 - `struct foo *f = (struct foo *) malloc(sizeof(struct foo) * 10);`
...
`free(f);`
...
`free(f);`
 - May cause either a use after free (because something else called `malloc()` and got that address) or corrupt `malloc`'s data (because you are no longer freeing a pointer called by `malloc`)

And Valgrind...

- Valgrind slows down your program by an order of magnitude, but...
 - It adds a tons of checks designed to catch most (but not all) memory errors
- Memory leaks
- Misuse of free
- Writing over the end of arrays
- You ***must*** run your program in Valgrind before you ask for debugging help from a TA!
- Tools like Valgrind are absolutely essential for debugging C code

And In Conclusion, ...

- C has three main memory segments in which to allocate data:
 - Static Data: Variables outside functions
 - Stack: Variables local to function
 - Heap: Objects explicitly malloc-ed/free-d.
- Heap data is biggest source of bugs in C code